

Error Detection and Corrective Feedback using Anomaly Detection in Long Jump Athletes

Pradyumn Vikram

July 8, 2025

Contents

Abstract

This report provides a descriptive analysis of the process undertaken in building Project Kitty - An analytics module developed to detect and provide corrective feedback to long jump athletes using video footage.

1 Choice of Sport and Rationale

Long jump has long been a beloved athletic sport, with players from 60+ countries taking part in the Olympics for long jumping. Keeping aside its popularity, it is also a highly technical sport, requiring technique and skill. Correction in posture as well as general approach to the sport if optimised, can help an athlete improve his/her performance by a large margin. We have focused on various factors playing a vital role in the distance covered by an athlete during a long jump such as -

- Hip flexion
- Knee flexion
- Trunk angle
- Run-up velocity

The project aims to provide descriptive feedback on the above mentioned factors to an athlete. The factors to evaluate have been chosen by referring to online sources on biomechanics.

2 Data Collection and Extraction of Long Jump Clips from Extended Videos

I have used the YouTubeDL and youtubearch module to extract feed from the Olympic games of four different years - Beijing, Rio, Tokyo and Paris. The extracted videos being very large in size for preprocessing have been split into smaller segments of approximately 50mb in size using the ffmpeg module, for ease of processing and extracting segments. I had to change the codec of videos before splitting them due to a system error, due to which they were transferred to the altered `odecdirectory`.

For extraction of segments of athletes performing long jumps, I first manually cropped a singular reference video (Take-off, Flight and Landing). The goal is now to find timestamps in the game feed videos representing

First, I tried to extract the perceptual hash of the reference video, and compare it to fixed segments (frames) in the videos where I want to extract clips from. However this approach did not work as -

- The video files differed too much, due to differences in the track along with visual differences in various aspects, such as jersey color etc.
- The visual summary generated (collage of images) could not capture the temporal data, as well as the features needed to extract the required frames.

Second, I used the ResNet3D-18 pretrained CNN (trained on the Kinetics 400 dataset - comprising of data involving human action classes).

- Each video snippet was segmented in 16 frames - each of which was then standardised with the mean and standard deviation of the kinetics 400 dataset
- These frames were then stacked to form a tensor of shape [channel, frames, width, height], and then reshaped to [1, channel, frames, width, height] (this is required for input to the ResNet model)

- This was then passed into the stem layer and 4 residual blocks of the pre-trained residual network to extract features from temporal data
- Now we extract the mean of the frames, width and height column; flattening the array furthermore, to represent the entire sequence in an array of 1 dimension (a vector)
- This was then compared using cosine similarity with the extracted 1 dimensional vector of the reference video
- Those scoring above a set threshold were flagged and then later joined with other segments who overlap to get a complete video clip
- These videos are then clipped according to the extracted timestamps, giving very accurate results

This approach, using the Residual Network returned very accurate clips of athletes performing long jumps, showing that the network managed to capture the temporal features very well. However not entirely immune to error, I had to iterate through all extracted clips and remove 15 out of 240 clips which did not contain the frames required

3 Phase Segmentation of Long Jump Clips

Now that we have extracted clips of athletes performing long jumps and eliminated a majority of the noise in the form of commentator box frames, audience etc. We now attempt to divide the clip into phases for further analysis, namely - (Run up, Take off, Flight and Landing). Since I had no labelled data, clustering seemed to be the most appropriate approach. I used mediapipe to extract features from each frame, which were then passed into the K-Means clustering algorithm to label each frame with an ID from 0 to 4.

The features extracted namely -

- Hip Velocity in X
- Hip Velocity in Y

- Hip Y coordinate
- Trunk angle with the Z-axis
- Hip-knee angle
- Right knee angle
- Left Knee angle
- Lowest foot coordinate

were chosen so via online research and trial and error, the evolution of these features through-out the frames can be seen in Figure 1

K-Means clustering was then employed after standardisation over all the features giving suitable output with some error. The clusters were then smoothened by using the fact that one phase when completed cannot be repeated again, see Figure 2

This clustering after employed on each frame, then enabled me to divide the phases into four datasets - corresponding to each phase. These can now be used for anomaly detection and correction by engineering appropriate features.

4 Anomaly Detection and Corrective Prediction

For anomaly detection ie isolating out errors and providing descriptive corrective feedback, I considered using Isolation Forests or One Class SVMs. Isolation Forests was ruled out since it would be difficult to identify which feature caused the anomaly and isolated the sample from the normal. I engineered 2 features each for each phase based on online research and prioritising important factors for good technique. Each data frame once standardised was then fit to a One Class SVM. The decision function plotted can be reviewed in Figure 3

For providing descriptive feedback, I first identify outliers (any point not in the contour), then I evaluate it's x and y distance from the center of the contour. Since x and y represent the 2 features based on which the contour was made, whichever distance is greater tell me, what was




Figure 1: Evolution of Features over Frames




Figure 2: Clustered Frames - Before vs After smoothing

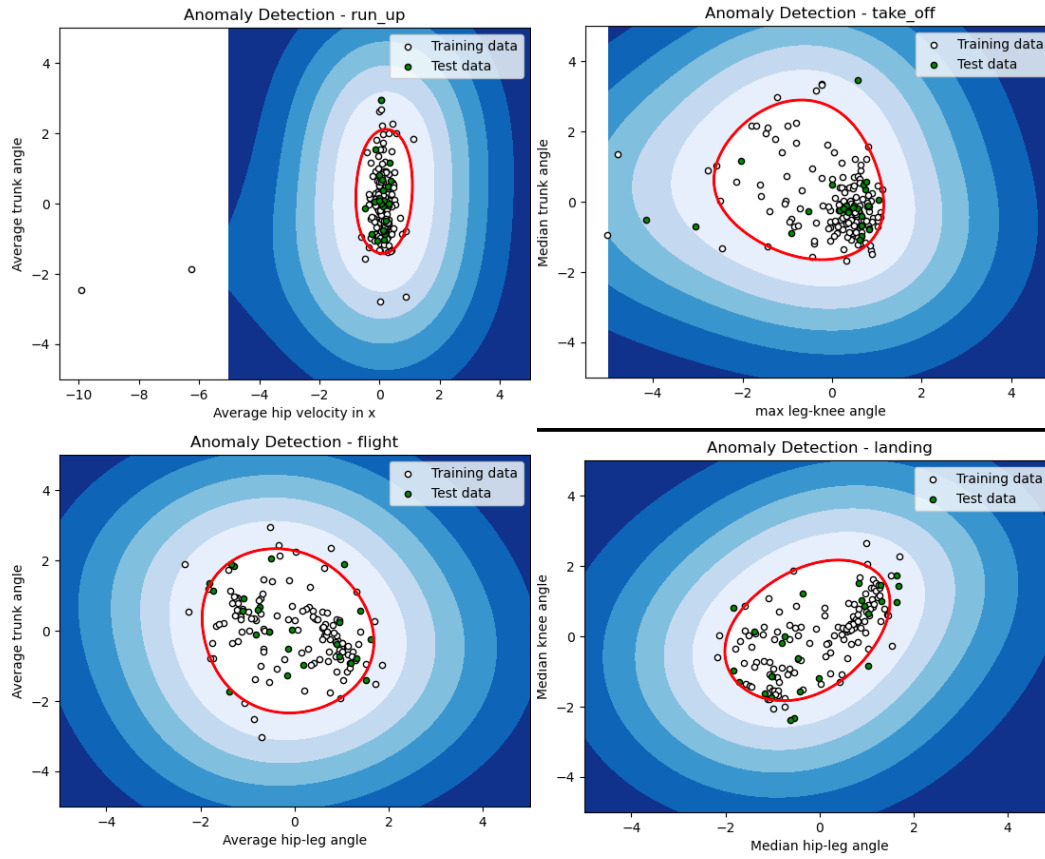


Figure 3: Decision Boundaries - Showing model fit and performance

the major contributor to the anomaly - something I cannot do with Isolation Forests. Then the feature contributing most to the anomaly is then acted upon to provide feedback to the athlete in four phases.

Proceeding to build the analytics module, I took an unseen data sample (Trimmed video of an athlete performing long jump from the Asian Games - refer to misc/), in which the model provided descriptive feedback. The interface for the module and the output can be seen in Figure 4

```
(base) pradyumnvikram@azidozide:~/projects/ProjectKitty$ python -u "/home/pradyumnvikram/projects/ProjectKitty/analytics module.py"
2025-07-07 22:48:29.177706: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
2025-07-07 22:48:29.188249: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1751908709.200170 280950 cuda.dnn.cc:8310] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered
E0000 00:00:1751908709.203751 280950 cuda.blas.cc:1418] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered
2025-07-07 22:48:29.216346: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 AVX VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
[LOG]: Enter path to cropped video file: misc/trimmed_test.webm
[LOG]: Loading models
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
I0000 00:00:1751908719.331545 280950 gl_context_egl.cc:85] Successfully initialized EGL. Major : 1 Minor: 5
I0000 00:00:1751908719.337399 281141 gl_context.cc:369] GL version: 3.2 (OpenGL ES 3.2 Mesa 23.2.1-lubuntu3.1-22.04.3), renderer: Mesa Intel(R) Graphics (RPL-S)
[LOG]: Extracting features
INFO: Created TensorFlow Lite XNNPACK delegate for CPU.
W0000 00:00:1751908719.414773 281100 inference_feedback_manager.cc:114] Feedback manager requires a model with a single signature inference. Disabling support for feedback tensors.
W0000 00:00:1751908719.447002 281118 inference_feedback_manager.cc:114] Feedback manager requires a model with a single signature inference. Disabling support for feedback tensors.
W0000 00:00:1751908719.462139 281118 landmark_projection_calculator.cc:186] Using NORM_RECT without IMAGE_DIMENSIONS is only supported for the square ROI. Provide IMAGE_DIMENSIONS or use PROJECTION_MATRIX.
[ANOMALY - Run Up]: Increase Speed
[ANOMALY - Take Off]: Increase knee flexion
[ANOMALY - Flight]: Maintain greater hip-leg angle
[ANOMALY - Landing]: Push chest towards legs
```

Figure 4: Analytics module - UI and Test output

5 Summary and Drawbacks

We have developed a module which can evaluate long jumps and provide corrective feedback.

- The pipeline is still sometimes struggling to cluster some video clips into various phases
- Since I have only unlabelled data, it is difficult to analyse the performance of the analytics module (The visualisation is my only source of verification)
- More features can be added for descriptive feedback - seeing how currently there are only 2 per phase